

```
In [ ]: import torch
import torch.nn as nn
import torch.optim as optim
from torchvision.datasets import CIFAR10
from torch.utils.data import DataLoader, Subset
import torchvision.transforms as transforms
from tqdm import tqdm
import matplotlib.pyplot as plt
import random

# Define the transformations
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)) # Normalize to [-1, 1]
])

# Normalization Importance:
# Improves model training stability and performance.
# Speeds up convergence by centering inputs around zero.
# Works well with activation functions like tanh and ReLU.

# Load the CIFAR-10 dataset
train_dataset = CIFAR10(root='./data', train=True, download=True, transform=transform)
test_dataset = CIFAR10(root='./data', train=False, download=True, transform=transform)

# Subset the dataset to 2500 examples
train_dataset = Subset(train_dataset, range(2500))
test_dataset = Subset(test_dataset, range(500))

# Visualize 5 random images
f, grid = plt.subplots(1, 5, figsize=(15, 5))
for i in range(5):
    idx = random.randint(0, 2499)
    # Convert tensor back to image for display
    image, label = train_dataset[idx]
    image = image.permute(1, 2, 0) # Convert from [C, H, W] to [H, W, C]
    image = (image * 0.5) + 0.5 # Unnormalize to [0, 1]
    grid[i].imshow(image.numpy())
    grid[i].axis('off')
    grid[i].set_title(f"Label: {label}")
```

```
plt.tight_layout()
plt.show()
```

Create DataLoaders

```
train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=64, shuffle=False)
```

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to ./data/cifar-10-python.tar.gz

100%|██████████| 170M/170M [00:13<00:00, 12.5MB/s]

Extracting ./data/cifar-10-python.tar.gz to ./data

Files already downloaded and verified



```
In [ ]: class LeNet5(nn.Module):
    def __init__(self):
        super(LeNet5, self).__init__()
        # First convolutional layer: 6 filters, kernel size 5x5
        self.conv1 = nn.Conv2d(3, 6, kernel_size=5, stride=1, padding=0)
        self.pool1 = nn.AvgPool2d(kernel_size=2, stride=2) # First pooling layer

        # Second convolutional layer: 16 filters, kernel size 5x5
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5, stride=1, padding=0)
        self.pool2 = nn.AvgPool2d(kernel_size=2, stride=2) # Second pooling layer

        # Fully connected layers
        self.fc1 = nn.Linear(16 * 5 * 5, 120) # Input: flattened feature maps, Output: 120 units
        self.fc2 = nn.Linear(120, 84) # Output: 84 units
        self.fc3 = nn.Linear(84, 10) # Output: 10 units (CIFAR-10 classes)

    def forward(self, x):
        # Forward pass through convolutional and pooling layers
```

```
x = torch.relu(self.conv1(x)) # Conv1 + ReLU
x = self.pool1(x) # Pooling1
x = torch.relu(self.conv2(x)) # Conv2 + ReLU
x = self.pool2(x) # Pooling2

# Flatten the tensor
x = x.view(-1, 16 * 5 * 5)

# Fully connected layers
x = torch.relu(self.fc1(x)) # FC1 + ReLU
x = torch.relu(self.fc2(x)) # FC2 + ReLU
x = self.fc3(x) # Output layer
return x

# Initialize the model, loss function, and optimizer
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = LeNet5().to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

# Store metrics
train_losses = []
val_losses = []
train_accuracies = []
val_accuracies = []

# Training loop with metrics logging
epochs = 10
for epoch in range(epochs):
    # Training Phase
    model.train()
    running_loss = 0.0
    correct_train = 0
    total_train = 0

    for images, labels in train_loader:
        images, labels = images.to(device), labels.to(device)

        # Forward pass
        outputs = model(images)
        loss = criterion(outputs, labels)
```

```
# Backward pass and optimization
optimizer.zero_grad()
loss.backward()
optimizer.step()

# Update metrics
running_loss += loss.item()
_, predicted = torch.max(outputs, 1)
correct_train += (predicted == labels).sum().item()
total_train += labels.size(0)

train_accuracy = correct_train / total_train
train_loss = running_loss / len(train_loader)
train_losses.append(train_loss)
train_accuracies.append(train_accuracy)

# Validation Phase
model.eval()
val_loss = 0.0
correct_val = 0
total_val = 0
with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        loss = criterion(outputs, labels)

        val_loss += loss.item()
        _, predicted = torch.max(outputs, 1)
        correct_val += (predicted == labels).sum().item()
        total_val += labels.size(0)

val_accuracy = correct_val / total_val
val_loss = val_loss / len(test_loader)
val_losses.append(val_loss)
val_accuracies.append(val_accuracy)

# Print epoch results
print(f"Epoch {epoch + 1}/{epochs}")
print(f"Train Loss: {train_loss:.4f}, Train Accuracy: {train_accuracy:.4f}, "
      f"Val Loss: {val_loss:.4f}, Val Accuracy: {val_accuracy:.4f}")
```

```

Epoch 1/10
Train Loss: 2.2317, Train Accuracy: 0.1720, Val Loss: 2.1511, Val Accuracy: 0.2280
Epoch 2/10
Train Loss: 2.0610, Train Accuracy: 0.2616, Val Loss: 2.0200, Val Accuracy: 0.2760
Epoch 3/10
Train Loss: 1.9774, Train Accuracy: 0.2816, Val Loss: 1.9612, Val Accuracy: 0.2960
Epoch 4/10
Train Loss: 1.8693, Train Accuracy: 0.3176, Val Loss: 1.8735, Val Accuracy: 0.3200
Epoch 5/10
Train Loss: 1.8185, Train Accuracy: 0.3420, Val Loss: 1.8527, Val Accuracy: 0.3300
Epoch 6/10
Train Loss: 1.7795, Train Accuracy: 0.3592, Val Loss: 1.8428, Val Accuracy: 0.3580
Epoch 7/10
Train Loss: 1.7594, Train Accuracy: 0.3592, Val Loss: 1.7572, Val Accuracy: 0.3560
Epoch 8/10
Train Loss: 1.6570, Train Accuracy: 0.3832, Val Loss: 1.7467, Val Accuracy: 0.3700
Epoch 9/10
Train Loss: 1.6256, Train Accuracy: 0.4140, Val Loss: 1.7266, Val Accuracy: 0.3980
Epoch 10/10
Train Loss: 1.5556, Train Accuracy: 0.4348, Val Loss: 1.7090, Val Accuracy: 0.4020

```

```

In [ ]: import matplotlib.pyplot as plt

# Plot Training and Validation Accuracy
plt.figure(figsize=(12, 5))

# Loss plot
plt.subplot(1, 2, 1)
plt.plot(range(1, epochs + 1), train_losses, label="Train Loss")
plt.plot(range(1, epochs + 1), val_losses, label="Validation Loss")
plt.title("Loss over Epochs")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()

# Accuracy plot
plt.subplot(1, 2, 2)
plt.plot(range(1, epochs + 1), train_accuracies, label="Train Accuracy")
plt.plot(range(1, epochs + 1), val_accuracies, label="Validation Accuracy")
plt.title("Accuracy over Epochs")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")

```

```
plt.legend()

plt.tight_layout()
plt.show()
```



```
In [ ]: # Training with L1 Regularization
def train_with_l1_regularization(model, train_loader, test_loader, device, lambda_l1=1e-5, epochs=10):
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.Adam(model.parameters(), lr=0.001)
    train_losses, val_losses = [], []
    train_accuracies, val_accuracies = [], []

    for epoch in range(epochs):
        model.train()
        running_loss = 0.0
        correct_train, total_train = 0, 0

        for images, labels in train_loader:
            images, labels = images.to(device), labels.to(device)
            outputs = model(images)
```

```
loss = criterion(outputs, labels)

# Add L1 regularization
l1_loss = sum(param.abs().sum() for param in model.parameters())
loss += lambda_l1 * l1_loss

optimizer.zero_grad()
loss.backward()
optimizer.step()

running_loss += loss.item()
_, predicted = torch.max(outputs, 1)
correct_train += (predicted == labels).sum().item()
total_train += labels.size(0)

train_loss = running_loss / len(train_loader)
train_accuracy = correct_train / total_train
train_losses.append(train_loss)
train_accuracies.append(train_accuracy)

# Validation phase
model.eval()
val_loss, correct_val, total_val = 0.0, 0, 0
with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        loss = criterion(outputs, labels)
        val_loss += loss.item()
        _, predicted = torch.max(outputs, 1)
        correct_val += (predicted == labels).sum().item()
        total_val += labels.size(0)

val_loss /= len(test_loader)
val_accuracy = correct_val / total_val
val_losses.append(val_loss)
val_accuracies.append(val_accuracy)

print(f"Epoch {epoch + 1}/{epochs}")
print(f"Train Loss (L1): {train_loss:.4f}, Train Accuracy: {train_accuracy:.4f}, "
      f"Val Loss: {val_loss:.4f}, Val Accuracy: {val_accuracy:.4f}")
```

```
    return train_losses, train_accuracies, val_losses, val_accuracies

# Training with L2 Regularization
def train_with_l2_regularization(model, train_loader, test_loader, device, lambda_l2=1e-5, epochs=10):
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.Adam(model.parameters(), lr=0.001, weight_decay=lambda_l2)
    train_losses, val_losses = [], []
    train_accuracies, val_accuracies = [], []

    for epoch in range(epochs):
        model.train()
        running_loss = 0.0
        correct_train, total_train = 0, 0

        for images, labels in train_loader:
            images, labels = images.to(device), labels.to(device)
            outputs = model(images)
            loss = criterion(outputs, labels)
            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

            running_loss += loss.item()
            _, predicted = torch.max(outputs, 1)
            correct_train += (predicted == labels).sum().item()
            total_train += labels.size(0)

        train_loss = running_loss / len(train_loader)
        train_accuracy = correct_train / total_train
        train_losses.append(train_loss)
        train_accuracies.append(train_accuracy)

    # Validation phase
    model.eval()
    val_loss, correct_val, total_val = 0.0, 0, 0
    with torch.no_grad():
        for images, labels in test_loader:
            images, labels = images.to(device), labels.to(device)
            outputs = model(images)
            loss = criterion(outputs, labels)
            val_loss += loss.item()
            _, predicted = torch.max(outputs, 1)
```



```
        correct_val += (predicted == labels).sum().item()
        total_val += labels.size(0)

    val_loss /= len(test_loader)
    val_accuracy = correct_val / total_val
    val_losses.append(val_loss)
    val_accuracies.append(val_accuracy)

    print(f"Epoch {epoch + 1}/{epochs}")
    print(f"Train Loss (L2): {train_loss:.4f}, Train Accuracy: {train_accuracy:.4f}, "
          f"Val Loss: {val_loss:.4f}, Val Accuracy: {val_accuracy:.4f}")

    return train_losses, train_accuracies, val_losses, val_accuracies

# Plotting function
def plot_metrics(train_losses, val_losses, train_accuracies, val_accuracies, title):
    plt.figure(figsize=(12, 5))

    # Plot losses
    plt.subplot(1, 2, 1)
    plt.plot(train_losses, label="Train Loss")
    plt.plot(val_losses, label="Validation Loss")
    plt.xlabel("Epochs")
    plt.ylabel("Loss")
    plt.title(f"{title} Loss")
    plt.legend()

    # Plot accuracies
    plt.subplot(1, 2, 2)
    plt.plot(train_accuracies, label="Train Accuracy")
    plt.plot(val_accuracies, label="Validation Accuracy")
    plt.xlabel("Epochs")
    plt.ylabel("Accuracy")
    plt.title(f"{title} Accuracy")
    plt.legend()

    plt.tight_layout()
    plt.show()

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model_l1 = LeNet5().to(device)
train_losses, train_accuracies, val_losses, val_accuracies = train_with_l1_regularization(
```

```
model_l1, train_loader, test_loader, device, lambda_l1=1e-5, epochs=10)
plot_metrics(train_losses, val_losses, train_accuracies, val_accuracies, "L1 Regularization")

model_l2 = LeNet5().to(device)
train_losses, train_accuracies, val_losses, val_accuracies = train_with_l2_regularization(
    model_l2, train_loader, test_loader, device, lambda_l2=1e-5, epochs=10)
plot_metrics(train_losses, val_losses, train_accuracies, val_accuracies, "L2 Regularization")
```

Epoch 1/10

Train Loss (L1): 2.2464, Train Accuracy: 0.1488, Val Loss: 2.1146, Val Accuracy: 0.2320

Epoch 2/10

Train Loss (L1): 2.0696, Train Accuracy: 0.2472, Val Loss: 1.9714, Val Accuracy: 0.2760

Epoch 3/10

Train Loss (L1): 1.9283, Train Accuracy: 0.2976, Val Loss: 1.8833, Val Accuracy: 0.2760

Epoch 4/10

Train Loss (L1): 1.8411, Train Accuracy: 0.3404, Val Loss: 1.8941, Val Accuracy: 0.3160

Epoch 5/10

Train Loss (L1): 1.7991, Train Accuracy: 0.3468, Val Loss: 1.7645, Val Accuracy: 0.3660

Epoch 6/10

Train Loss (L1): 1.7284, Train Accuracy: 0.3744, Val Loss: 1.7949, Val Accuracy: 0.3340

Epoch 7/10

Train Loss (L1): 1.6828, Train Accuracy: 0.3892, Val Loss: 1.7618, Val Accuracy: 0.3400

Epoch 8/10


Train Loss (L1): 1.6391, Train Accuracy: 0.4084, Val Loss: 1.7302, Val Accuracy: 0.3620

Epoch 9/10

Train Loss (L1): 1.5825, Train Accuracy: 0.4328, Val Loss: 1.7321, Val Accuracy: 0.3740

Epoch 10/10

Train Loss (L1): 1.5687, Train Accuracy: 0.4460, Val Loss: 1.6845, Val Accuracy: 0.3480



Epoch 1/10

Train Loss (L2): 2.2142, Train Accuracy: 0.1636, Val Loss: 2.0369, Val Accuracy: 0.2280

Epoch 2/10

Train Loss (L2): 2.0027, Train Accuracy: 0.2580, Val Loss: 1.9607, Val Accuracy: 0.2540

Epoch 3/10

Train Loss (L2): 1.9328, Train Accuracy: 0.3016, Val Loss: 1.9534, Val Accuracy: 0.3020

Epoch 4/10

Train Loss (L2): 1.8728, Train Accuracy: 0.3216, Val Loss: 1.8714, Val Accuracy: 0.3340

Epoch 5/10

Train Loss (L2): 1.7994, Train Accuracy: 0.3456, Val Loss: 1.7971, Val Accuracy: 0.3700

Epoch 6/10

Train Loss (L2): 1.7379, Train Accuracy: 0.3748, Val Loss: 1.8040, Val Accuracy: 0.3360

Epoch 7/10

Train Loss (L2): 1.6978, Train Accuracy: 0.3920, Val Loss: 1.7420, Val Accuracy: 0.3500

Epoch 8/10


Train Loss (L2): 1.6505, Train Accuracy: 0.3916, Val Loss: 1.7224, Val Accuracy: 0.3580

Epoch 9/10

Train Loss (L2): 1.6266, Train Accuracy: 0.4180, Val Loss: 1.7685, Val Accuracy: 0.3680

Epoch 10/10

Train Loss (L2): 1.5734, Train Accuracy: 0.4264, Val Loss: 1.7060, Val Accuracy: 0.3800



```
In [ ]: class LeNet5Flexible(nn.Module):
    def __init__(self, activation_function):
        super(LeNet5Flexible, self).__init__()
        # First convolutional layer: 6 filters, kernel size 5x5
        self.conv1 = nn.Conv2d(3, 6, kernel_size=5, stride=1, padding=0)
        self.pool1 = nn.AvgPool2d(kernel_size=2, stride=2) # First pooling layer

        # Second convolutional layer: 16 filters, kernel size 5x5
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5, stride=1, padding=0)
        self.pool2 = nn.AvgPool2d(kernel_size=2, stride=2) # Second pooling layer

        # Fully connected layers
        self.fc1 = nn.Linear(16 * 5 * 5, 120) # Input: flattened feature maps, Output: 120 units
        self.fc2 = nn.Linear(120, 84) # Output: 84 units
        self.fc3 = nn.Linear(84, 10) # Output: 10 units (CIFAR-10 classes)

        # Dynamic activation function
        self.activation = activation_function

    def forward(self, x):
```

```

# Forward pass through convolutional and pooling layers
x = self.activation(self.conv1(x)) # Conv1 + Activation
x = self.pool1(x) # Pooling1
x = self.activation(self.conv2(x)) # Conv2 + Activation
x = self.pool2(x) # Pooling2

# Flatten the tensor
x = x.view(-1, 16 * 5 * 5)

# Fully connected layers
x = self.activation(self.fc1(x)) # FC1 + Activation
x = self.activation(self.fc2(x)) # FC2 + Activation
x = self.fc3(x) # Output layer (no activation here, handled by loss or softmax later)
return x

```

```

In [ ]: # Define the device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Define the activation functions
activation_functions = {
    "Tanh": torch.tanh,
    "Leaky ReLU": nn.LeakyReLU(),
    "Softmax": lambda x: nn.functional.softmax(x, dim=1), # Softmax applied layer-wise
}

def train_and_collect_metrics(model, train_loader, test_loader, device, epochs=10):
    # Define the loss criterion and optimizer
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.Adam(model.parameters(), lr=0.001)

    # Lists to store metrics
    train_losses, val_losses = [], []
    train_accuracies, val_accuracies = [], []

    for epoch in range(epochs):
        # Training Phase
        model.train()
        running_loss = 0.0
        correct_train, total_train = 0, 0

        for images, labels in train_loader:
            images, labels = images.to(device), labels.to(device)

```

```
# Forward pass
outputs = model(images)
loss = criterion(outputs, labels)

# Backward pass and optimization
optimizer.zero_grad()
loss.backward()
optimizer.step()

# Accumulate metrics
running_loss += loss.item()
_, predicted = torch.max(outputs, 1)
correct_train += (predicted == labels).sum().item()
total_train += labels.size(0)

train_losses.append(running_loss / len(train_loader))
train_accuracies.append(correct_train / total_train)

# Validation Phase
model.eval()
val_loss, correct_val, total_val = 0.0, 0, 0
with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        loss = criterion(outputs, labels)
        val_loss += loss.item()
        _, predicted = torch.max(outputs, 1)
        correct_val += (predicted == labels).sum().item()
        total_val += labels.size(0)

val_losses.append(val_loss / len(test_loader))
val_accuracies.append(correct_val / total_val)

# Print progress
print(f"Epoch {epoch + 1}/{epochs}")
print(f"Train Loss: {train_losses[-1]:.4f}, Train Accuracy: {train_accuracies[-1]:.4f}, "
      f"Val Loss: {val_losses[-1]:.4f}, Val Accuracy: {val_accuracies[-1]:.4f}")

return train_losses, train_accuracies, val_losses, val_accuracies
```

```
# Train the model for each activation function
results = {}
for name, activation_function in activation_functions.items():
    print(f"\nTraining with {name} activation function\n")

    # Initialize the model
    model = LeNet5Flexible(activation_function=activation_function).to(device)

    # Train the model (replace train_loader and test_loader with your actual data loaders)
    train_losses, train_accuracies, val_losses, val_accuracies = train_and_collect_metrics(
        model, train_loader, test_loader, device, epochs=10
    )

    # Save the results for analysis
    results[name] = {
        "train_losses": train_losses,
        "val_losses": val_losses,
        "train_accuracies": train_accuracies,
        "val_accuracies": val_accuracies,
    }

    # Plot the metrics for this activation function
    plot_metrics(
        train_losses, val_losses, train_accuracies, val_accuracies, title=f"{name} Activation Function"
    )
```

Training with Tanh activation function

Epoch 1/10

Train Loss: 2.1041, Train Accuracy: 0.2464, Val Loss: 1.9705, Val Accuracy: 0.3180

Epoch 2/10

Train Loss: 1.9583, Train Accuracy: 0.2956, Val Loss: 1.9254, Val Accuracy: 0.3200

Epoch 3/10

Train Loss: 1.8996, Train Accuracy: 0.3288, Val Loss: 1.9139, Val Accuracy: 0.3020

Epoch 4/10

Train Loss: 1.8400, Train Accuracy: 0.3556, Val Loss: 1.8906, Val Accuracy: 0.3340

Epoch 5/10

Train Loss: 1.7713, Train Accuracy: 0.3704, Val Loss: 1.8743, Val Accuracy: 0.3200

Epoch 6/10

Train Loss: 1.7225, Train Accuracy: 0.3944, Val Loss: 1.8706, Val Accuracy: 0.3300

Epoch 7/10

Train Loss: 1.7004, Train Accuracy: 0.3924, Val Loss: 1.8713, Val Accuracy: 0.3180

Epoch 8/10


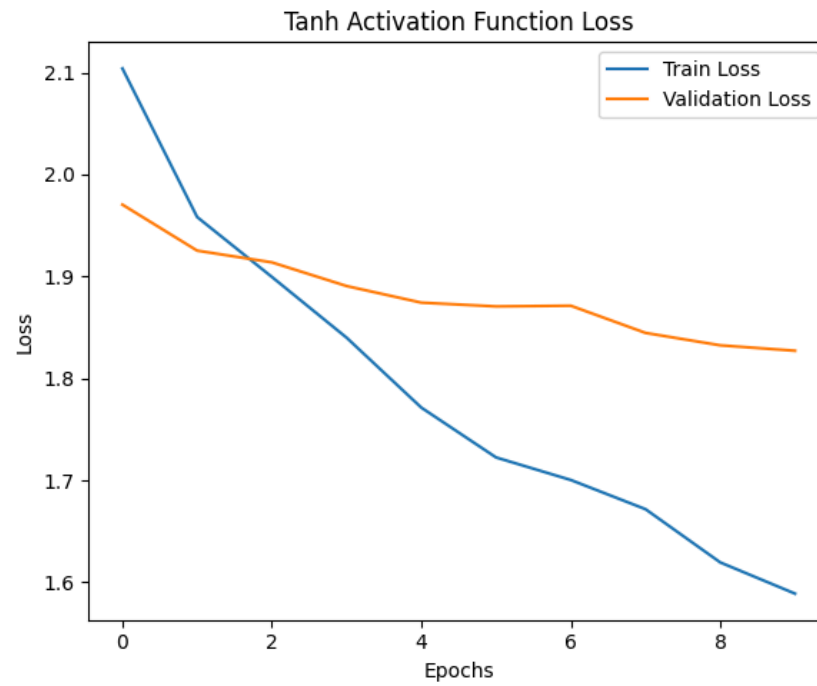
Train Loss: 1.6717, Train Accuracy: 0.4080, Val Loss: 1.8446, Val Accuracy: 0.3120

Epoch 9/10

Train Loss: 1.6195, Train Accuracy: 0.4304, Val Loss: 1.8325, Val Accuracy: 0.3380

Epoch 10/10

Train Loss: 1.5889, Train Accuracy: 0.4288, Val Loss: 1.8272, Val Accuracy: 0.3380



Training with Leaky ReLU activation function

Epoch 1/10

Train Loss: 2.2191, Train Accuracy: 0.1524, Val Loss: 2.0564, Val Accuracy: 0.2600

Epoch 2/10

Train Loss: 2.0086, Train Accuracy: 0.2604, Val Loss: 1.9741, Val Accuracy: 0.2740

Epoch 3/10

Train Loss: 1.9311, Train Accuracy: 0.3060, Val Loss: 1.9326, Val Accuracy: 0.2800

Epoch 4/10

Train Loss: 1.7998, Train Accuracy: 0.3528, Val Loss: 1.8268, Val Accuracy: 0.3200

Epoch 5/10

Train Loss: 1.7424, Train Accuracy: 0.3676, Val Loss: 1.8279, Val Accuracy: 0.3200

Epoch 6/10

Train Loss: 1.7197, Train Accuracy: 0.3660, Val Loss: 1.7717, Val Accuracy: 0.3520

Epoch 7/10

Train Loss: 1.6926, Train Accuracy: 0.3952, Val Loss: 1.7299, Val Accuracy: 0.3780

Epoch 8/10



Train Loss: 1.6263, Train Accuracy: 0.4100, Val Loss: 1.7603, Val Accuracy: 0.3600

Epoch 9/10

Train Loss: 1.5981, Train Accuracy: 0.4128, Val Loss: 1.6739, Val Accuracy: 0.3720

Epoch 10/10

Train Loss: 1.5765, Train Accuracy: 0.4112, Val Loss: 1.7426, Val Accuracy: 0.3800



Training with Softmax activation function

Epoch 1/10

Train Loss: 2.3047, Train Accuracy: 0.0944, Val Loss: 2.3010, Val Accuracy: 0.0960

Epoch 2/10

Train Loss: 2.3045, Train Accuracy: 0.0944, Val Loss: 2.3009, Val Accuracy: 0.0960

Epoch 3/10

Train Loss: 2.3040, Train Accuracy: 0.0920, Val Loss: 2.3008, Val Accuracy: 0.1080

Epoch 4/10

Train Loss: 2.3028, Train Accuracy: 0.1020, Val Loss: 2.3012, Val Accuracy: 0.1140

Epoch 5/10

Train Loss: 2.3026, Train Accuracy: 0.1140, Val Loss: 2.3006, Val Accuracy: 0.1960

Epoch 6/10

Train Loss: 2.3014, Train Accuracy: 0.1032, Val Loss: 2.2996, Val Accuracy: 0.1600

Epoch 7/10

Train Loss: 2.2999, Train Accuracy: 0.1688, Val Loss: 2.2978, Val Accuracy: 0.1580

Epoch 8/10

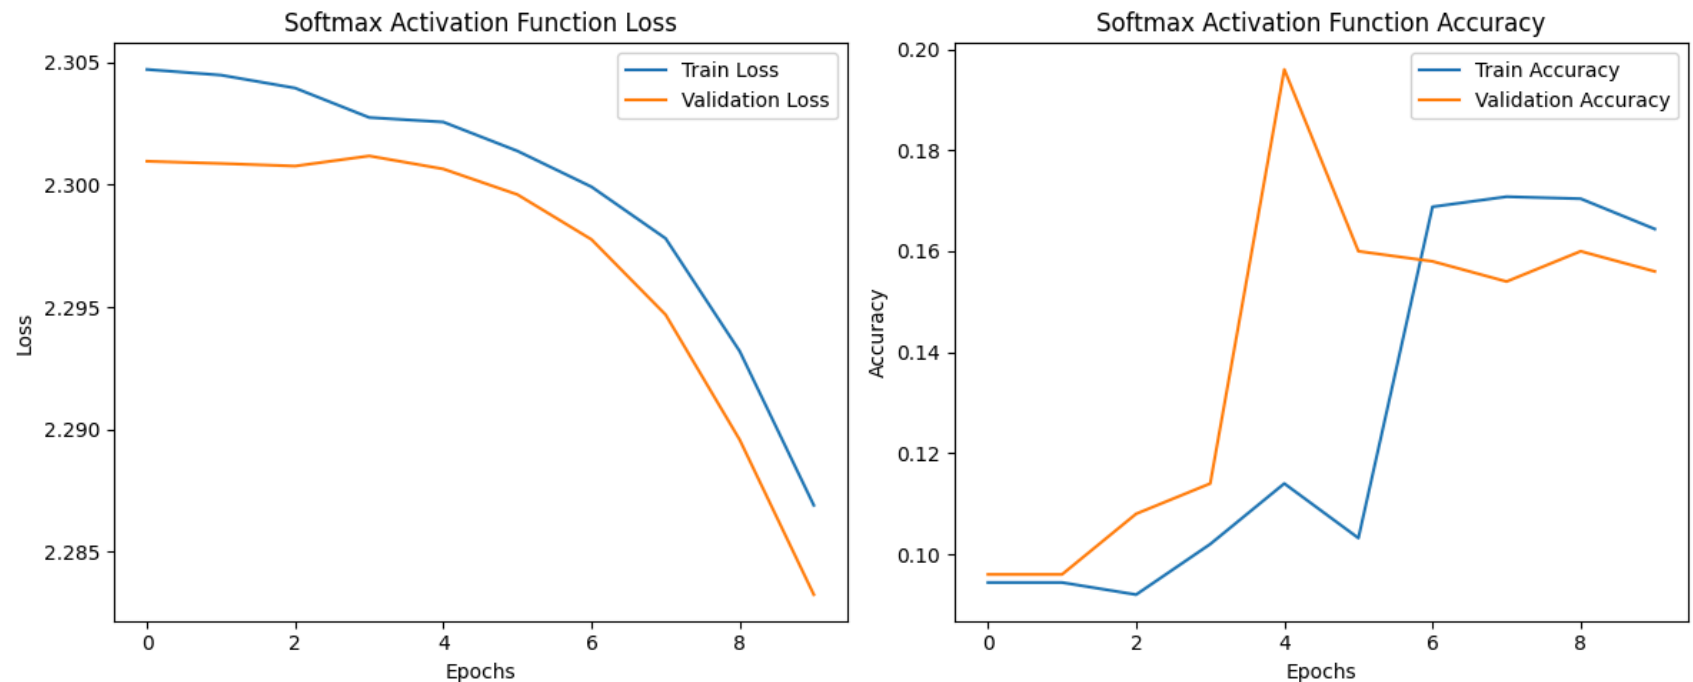
Train Loss: 2.2978, Train Accuracy: 0.1708, Val Loss: 2.2947, Val Accuracy: 0.1540

Epoch 9/10

Train Loss: 2.2932, Train Accuracy: 0.1704, Val Loss: 2.2896, Val Accuracy: 0.1600

Epoch 10/10

Train Loss: 2.2869, Train Accuracy: 0.1644, Val Loss: 2.2832, Val Accuracy: 0.1560



```
In [ ]: class LeNet5WithDropout(nn.Module):
    def __init__(self, dropout_rate=0.5):
        super(LeNet5WithDropout, self).__init__()
        # First convolutional layer: 6 filters, kernel size 5x5
        self.conv1 = nn.Conv2d(3, 6, kernel_size=5, stride=1, padding=0)
        self.pool1 = nn.AvgPool2d(kernel_size=2, stride=2) # First pooling layer

        # Second convolutional layer: 16 filters, kernel size 5x5
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5, stride=1, padding=0)
        self.pool2 = nn.AvgPool2d(kernel_size=2, stride=2) # Second pooling layer

        # Fully connected layers
        self.fc1 = nn.Linear(16 * 5 * 5, 120) # Input: flattened feature maps, Output: 120 units
        self.fc2 = nn.Linear(120, 84) # Output: 84 units
        self.fc3 = nn.Linear(84, 10) # Output: 10 units (CIFAR-10 classes)

        # Dropout
        self.dropout = nn.Dropout(p=dropout_rate)

    def forward(self, x):
```

```

    # Forward pass through convolutional and pooling layers
    x = torch.relu(self.conv1(x)) # Conv1 + ReLU
    x = self.pool1(x) # Pooling1
    x = torch.relu(self.conv2(x)) # Conv2 + ReLU
    x = self.pool2(x) # Pooling2

    # Flatten the tensor
    x = x.view(-1, 16 * 5 * 5)

    # Fully connected layers with dropout
    x = torch.relu(self.dropout(self.fc1(x))) # FC1 + ReLU + Dropout
    x = torch.relu(self.dropout(self.fc2(x))) # FC2 + ReLU + Dropout
    x = self.fc3(x) # Output layer
    return x
dropout_rates = [0.1, 0.2, 0.5]
results = {}

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

for rate in dropout_rates:
    print(f"\nTesting Dropout Rate: {rate}\n")
    model = LeNet5WithDropout(dropout_rate=rate).to(device)

    # Train the model (replace train_loader and test_loader with actual data loaders)
    train_losses, train_accuracies, val_losses, val_accuracies = train_and_collect_metrics(
        model, train_loader, test_loader, device, epochs=10
    )

    # Save results
    results[f"Dropout {rate}"] = {
        "train_losses": train_losses,
        "val_losses": val_losses,
        "train_accuracies": train_accuracies,
        "val_accuracies": val_accuracies,
    }

    # Plot metrics for this dropout rate
    plot_metrics(
        train_losses, val_losses, train_accuracies, val_accuracies, title=f"Dropout Rate {rate}"
    )

```

Testing Dropout Rate: 0.1

Epoch 1/10

Train Loss: 2.2089, Train Accuracy: 0.1736, Val Loss: 2.0105, Val Accuracy: 0.2860

Epoch 2/10

Train Loss: 2.0117, Train Accuracy: 0.2544, Val Loss: 1.9547, Val Accuracy: 0.2600

Epoch 3/10

Train Loss: 1.9372, Train Accuracy: 0.2888, Val Loss: 1.8768, Val Accuracy: 0.3140

Epoch 4/10

Train Loss: 1.8674, Train Accuracy: 0.3176, Val Loss: 1.8293, Val Accuracy: 0.3380

Epoch 5/10

Train Loss: 1.7901, Train Accuracy: 0.3400, Val Loss: 1.7846, Val Accuracy: 0.3300

Epoch 6/10

Train Loss: 1.7480, Train Accuracy: 0.3528, Val Loss: 1.7788, Val Accuracy: 0.3320

Epoch 7/10

Train Loss: 1.7232, Train Accuracy: 0.3656, Val Loss: 1.7542, Val Accuracy: 0.3260

Epoch 8/10

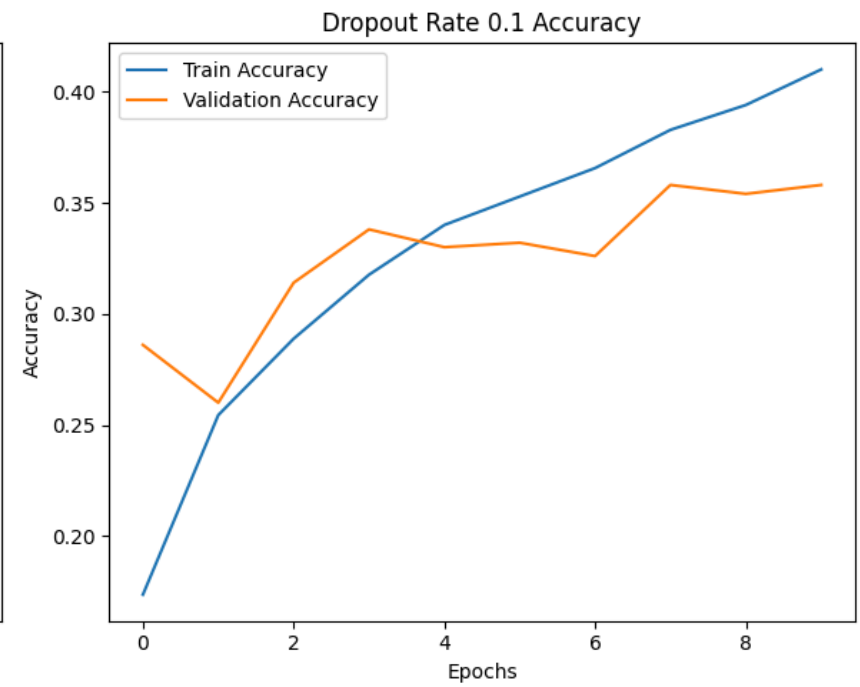
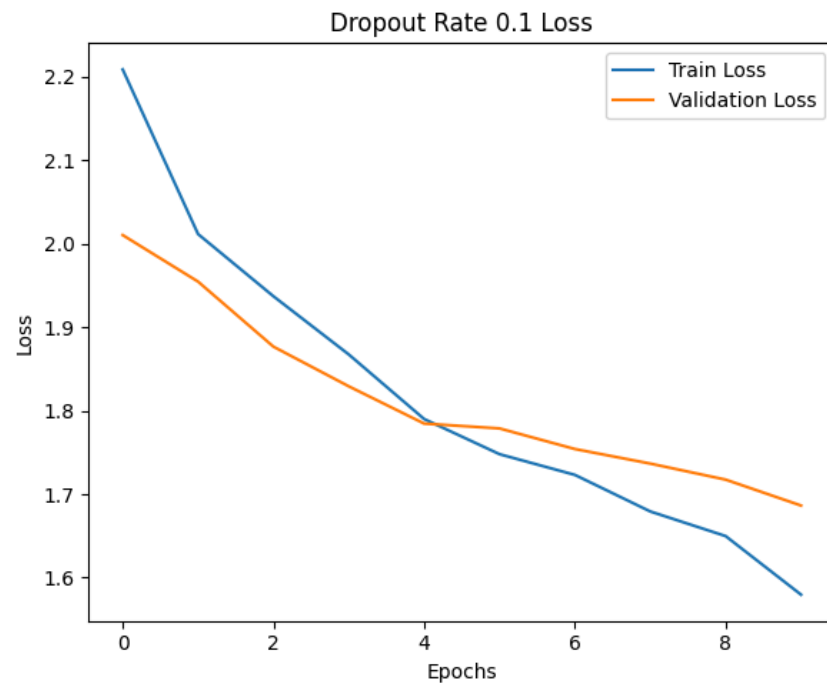
Train Loss: 1.6794, Train Accuracy: 0.3828, Val Loss: 1.7367, Val Accuracy: 0.3580

Epoch 9/10

Train Loss: 1.6499, Train Accuracy: 0.3940, Val Loss: 1.7175, Val Accuracy: 0.3540

Epoch 10/10

Train Loss: 1.5797, Train Accuracy: 0.4100, Val Loss: 1.6865, Val Accuracy: 0.3580



Testing Dropout Rate: 0.2

Epoch 1/10

Train Loss: 2.2339, Train Accuracy: 0.1500, Val Loss: 2.0960, Val Accuracy: 0.2100

Epoch 2/10

Train Loss: 2.0525, Train Accuracy: 0.2404, Val Loss: 1.9809, Val Accuracy: 0.2760

Epoch 3/10

Train Loss: 1.9434, Train Accuracy: 0.2888, Val Loss: 1.9050, Val Accuracy: 0.3320

Epoch 4/10

Train Loss: 1.8934, Train Accuracy: 0.3136, Val Loss: 1.9157, Val Accuracy: 0.2980

Epoch 5/10

Train Loss: 1.8447, Train Accuracy: 0.3360, Val Loss: 1.8612, Val Accuracy: 0.3040

Epoch 6/10

Train Loss: 1.7794, Train Accuracy: 0.3564, Val Loss: 1.8394, Val Accuracy: 0.3320

Epoch 7/10

Train Loss: 1.7896, Train Accuracy: 0.3476, Val Loss: 1.7793, Val Accuracy: 0.3600

Epoch 8/10



Train Loss: 1.7311, Train Accuracy: 0.3624, Val Loss: 1.7844, Val Accuracy: 0.3540

Epoch 9/10

Train Loss: 1.6596, Train Accuracy: 0.3928, Val Loss: 1.7641, Val Accuracy: 0.3500

Epoch 10/10

Train Loss: 1.6707, Train Accuracy: 0.3948, Val Loss: 1.7178, Val Accuracy: 0.3780



Testing Dropout Rate: 0.5

Epoch 1/10

Train Loss: 2.2720, Train Accuracy: 0.1348, Val Loss: 2.1377, Val Accuracy: 0.2200

Epoch 2/10

Train Loss: 2.1537, Train Accuracy: 0.2024, Val Loss: 2.0447, Val Accuracy: 0.2680

Epoch 3/10

Train Loss: 2.0750, Train Accuracy: 0.2176, Val Loss: 1.9889, Val Accuracy: 0.2740

Epoch 4/10

Train Loss: 2.0417, Train Accuracy: 0.2588, Val Loss: 1.9518, Val Accuracy: 0.2980

Epoch 5/10

Train Loss: 1.9885, Train Accuracy: 0.2448, Val Loss: 1.9230, Val Accuracy: 0.2980

Epoch 6/10

Train Loss: 1.9371, Train Accuracy: 0.2800, Val Loss: 1.8814, Val Accuracy: 0.3320

Epoch 7/10

Train Loss: 1.9158, Train Accuracy: 0.3128, Val Loss: 1.8674, Val Accuracy: 0.3480

Epoch 8/10


Train Loss: 1.9167, Train Accuracy: 0.3068, Val Loss: 1.8576, Val Accuracy: 0.3380

Epoch 9/10

Train Loss: 1.8979, Train Accuracy: 0.3100, Val Loss: 1.8375, Val Accuracy: 0.3460

Epoch 10/10

Train Loss: 1.8592, Train Accuracy: 0.3376, Val Loss: 1.8034, Val Accuracy: 0.3280



```
In [ ]: class LeNet5WithPooling(nn.Module):
    def __init__(self, pooling_type="avg"):
        super(LeNet5WithPooling, self).__init__()
        # First convolutional layer: 6 filters, kernel size 5x5
        self.conv1 = nn.Conv2d(3, 6, kernel_size=5, stride=1, padding=0)

        # Pooling layer
        if pooling_type == "avg":
            self.pool = nn.AvgPool2d(kernel_size=2, stride=2)
        elif pooling_type == "max":
            self.pool = nn.MaxPool2d(kernel_size=2, stride=2)
        elif pooling_type == "min":
            self.pool = nn.MaxPool2d(kernel_size=2, stride=2) # Min pooling simulated with negative MaxPool
            self.invert = True
        else:
            raise ValueError("Invalid pooling_type. Choose from ['avg', 'max', 'min'].")

        # Second convolutional layer: 16 filters, kernel size 5x5
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5, stride=1, padding=0)
```

```

    # Fully connected layers
    self.fc1 = nn.Linear(16 * 5 * 5, 120)
    self.fc2 = nn.Linear(120, 84)
    self.fc3 = nn.Linear(84, 10)

    # Store pooling type
    self.pooling_type = pooling_type

    def forward(self, x):
        # Apply pooling (invert input/output for min pooling)
        if self.pooling_type == "min":
            x = -self.pool(-torch.relu(self.conv1(x))) # Conv1 + Min Pool
            x = -self.pool(-torch.relu(self.conv2(x))) # Conv2 + Min Pool
        else:
            x = self.pool(torch.relu(self.conv1(x))) # Conv1 + Pool
            x = self.pool(torch.relu(self.conv2(x))) # Conv2 + Pool

        # Flatten the tensor
        x = x.view(-1, 16 * 5 * 5)

        # Fully connected layers
        x = torch.relu(self.fc1(x))
        x = torch.relu(self.fc2(x))
        x = self.fc3(x) # Output layer
        return x

```

```

In [ ]: pooling_types = ["avg", "max", "min"]
        results = {}

        device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

        for pooling_type in pooling_types:
            print(f"\nTesting {pooling_type} pooling\n")
            model = LeNet5WithPooling(pooling_type=pooling_type).to(device)

            # Train the model (replace train_loader and test_loader with actual data loaders)
            train_losses, train_accuracies, val_losses, val_accuracies = train_and_collect_metrics(
                model, train_loader, test_loader, device, epochs=10
            )

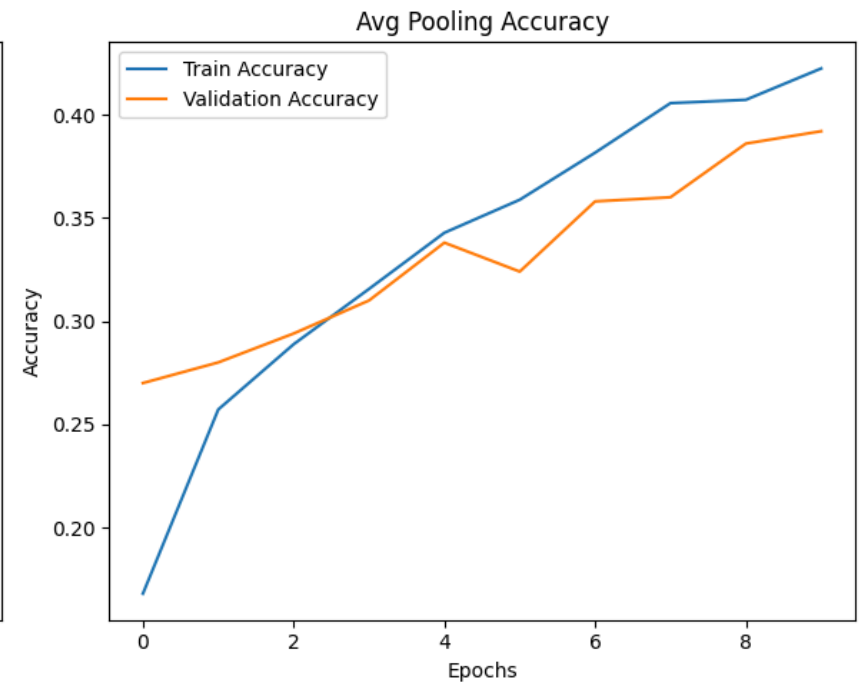
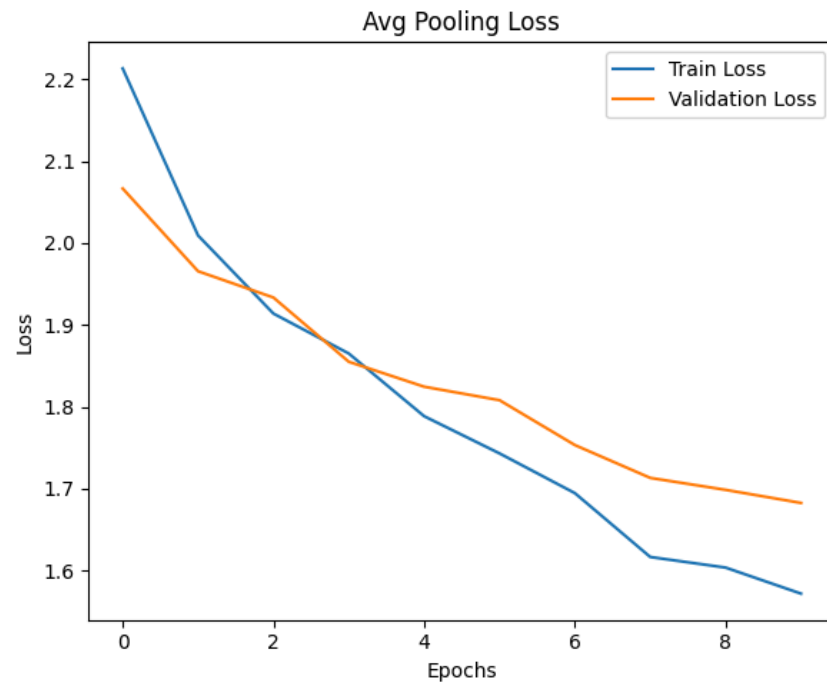
```

```
# Save results
results[f"{pooling_type} pooling"] = {
    "train_losses": train_losses,
    "val_losses": val_losses,
    "train_accuracies": train_accuracies,
    "val_accuracies": val_accuracies,
}

# Plot metrics for this pooling type
plot_metrics(
    train_losses, val_losses, train_accuracies, val_accuracies, title=f"{pooling_type.capitalize()} Po
)
```

Testing avg pooling

Epoch 1/10
Train Loss: 2.2133, Train Accuracy: 0.1680, Val Loss: 2.0666, Val Accuracy: 0.2700
Epoch 2/10
Train Loss: 2.0093, Train Accuracy: 0.2572, Val Loss: 1.9656, Val Accuracy: 0.2800
Epoch 3/10
Train Loss: 1.9139, Train Accuracy: 0.2888, Val Loss: 1.9335, Val Accuracy: 0.2940
Epoch 4/10
Train Loss: 1.8651, Train Accuracy: 0.3156, Val Loss: 1.8550, Val Accuracy: 0.3100
Epoch 5/10
Train Loss: 1.7885, Train Accuracy: 0.3428, Val Loss: 1.8245, Val Accuracy: 0.3380
Epoch 6/10
Train Loss: 1.7431, Train Accuracy: 0.3588, Val Loss: 1.8081, Val Accuracy: 0.3240
Epoch 7/10
Train Loss: 1.6945, Train Accuracy: 0.3816, Val Loss: 1.7533, Val Accuracy: 0.3580
Epoch 8/10
Train Loss: 1.6166, Train Accuracy: 0.4056, Val Loss: 1.7132, Val Accuracy: 0.3600
Epoch 9/10
Train Loss: 1.6036, Train Accuracy: 0.4072, Val Loss: 1.6986, Val Accuracy: 0.3860
Epoch 10/10
Train Loss: 1.5718, Train Accuracy: 0.4224, Val Loss: 1.6826, Val Accuracy: 0.3920



Testing max pooling

Epoch 1/10

Train Loss: 2.2311, Train Accuracy: 0.1508, Val Loss: 2.0998, Val Accuracy: 0.2600

Epoch 2/10

Train Loss: 1.9684, Train Accuracy: 0.2680, Val Loss: 1.9482, Val Accuracy: 0.2940

Epoch 3/10

Train Loss: 1.8720, Train Accuracy: 0.3024, Val Loss: 1.8684, Val Accuracy: 0.3260

Epoch 4/10

Train Loss: 1.7812, Train Accuracy: 0.3500, Val Loss: 1.8108, Val Accuracy: 0.3640

Epoch 5/10

Train Loss: 1.6904, Train Accuracy: 0.3844, Val Loss: 1.7716, Val Accuracy: 0.3380

Epoch 6/10

Train Loss: 1.6344, Train Accuracy: 0.3948, Val Loss: 1.7273, Val Accuracy: 0.3840

Epoch 7/10

Train Loss: 1.5633, Train Accuracy: 0.4276, Val Loss: 1.7361, Val Accuracy: 0.3560

Epoch 8/10


Train Loss: 1.5737, Train Accuracy: 0.4288, Val Loss: 1.7006, Val Accuracy: 0.3840

Epoch 9/10

Train Loss: 1.5352, Train Accuracy: 0.4584, Val Loss: 1.6722, Val Accuracy: 0.3740

Epoch 10/10

Train Loss: 1.4576, Train Accuracy: 0.4832, Val Loss: 1.6361, Val Accuracy: 0.3940



Testing min pooling

Epoch 1/10

Train Loss: 2.2327, Train Accuracy: 0.1544, Val Loss: 2.0737, Val Accuracy: 0.2540

Epoch 2/10

Train Loss: 2.0162, Train Accuracy: 0.2616, Val Loss: 1.9667, Val Accuracy: 0.3100

Epoch 3/10

Train Loss: 1.9277, Train Accuracy: 0.3032, Val Loss: 1.9044, Val Accuracy: 0.3320

Epoch 4/10

Train Loss: 1.8222, Train Accuracy: 0.3424, Val Loss: 1.8440, Val Accuracy: 0.2960

Epoch 5/10

Train Loss: 1.7470, Train Accuracy: 0.3556, Val Loss: 1.8214, Val Accuracy: 0.3020

Epoch 6/10

Train Loss: 1.6994, Train Accuracy: 0.3780, Val Loss: 1.8198, Val Accuracy: 0.3100

Epoch 7/10

Train Loss: 1.6704, Train Accuracy: 0.3928, Val Loss: 1.7611, Val Accuracy: 0.3620

Epoch 8/10

Train Loss: 1.6002, Train Accuracy: 0.4128, Val Loss: 1.7543, Val Accuracy: 0.3600

Epoch 9/10

Train Loss: 1.5439, Train Accuracy: 0.4364, Val Loss: 1.7487, Val Accuracy: 0.3640

Epoch 10/10

Train Loss: 1.5281, Train Accuracy: 0.4560, Val Loss: 1.7183, Val Accuracy: 0.3740

